



Degree Project in Computer Science and Engineering

Second cycle, 120 credits

FFS: A secure cloud-based steganographic filesystem through exploitation of online web services

Store your sensitive data in plain sight

GLENN OLSSON

FFS: A secure cloud-based steganographic filesystem through exploitation of online web services

Store your sensitive data in plain sight

GLENN OLSSON

Master's Programme, Computer Science, 120 credits
Date: April 24, 2022

Supervisors: Hamid Ghasemirahni, Zachory Peterson
Examiner: Gerald Quentin Maguire Jr

Host organization: Cal Poly
Swedish title: FFS: Ett säkert molnbaserat steganografiskt filsystem
genom utnyttjande av onlinebaserade webbtjänster
Swedish subtitle: Lagra din känsliga data öppet

Abstract

Today there are free online services that can be used to store files of arbitrary types and sizes, such as Google Drive. However, these services are often limited by a certain total storage size. The goal of this thesis is to create a filesystem that can store arbitrary amount and types of data, i.e. without any real limit to the total storage size. This is to be achieved by taking advantage of online webpages, such as Twitter, where text and files can be posted on free accounts with no apparent limit on storage size. The aim is to have a filesystem that behaves similar to any other filesystem but where the actual data is stored for free on various websites.

Keywords

Filesystem, Fejk FileSystem, Cloud-based filesystem, Steganographic filesystem

Sammanfattning

Sammanfattning på svenska

Nyckelord

Filsystem, Fejk FileSystem, Molnbaserat filsystem, Steganografiskt filsystem

Acknowledgments

Thanks to my mom, dad, and the rest of my family for their constant support

San Luis Obispo, California, USA April 2022

Glenn Olsson

Contents

1	Introduction	1
1.1	Project Overview	2
1.2	Problem	2
1.3	Purpose and motivation	3
1.4	Goals	4
1.5	Research Methodology	4
1.6	Delimitations	5
1.7	Structure of the thesis	6
2	Background	7
2.1	Filesystems and data storage	7
2.1.1	Unix filesystems	7
2.1.2	Distributed filesystems	9
2.1.3	Data storage and encoding	9
2.2	FUSE	11
2.3	Twitter	12
2.4	Threats	12
3	Related work	15
3.1	Steganography and deniable filesystems	15
3.2	Cryptography	16
3.3	Related filesystems	17
3.4	Filesystem benchmarking	18
3.5	Summary	19
4	Method	21
4.1	Development environment specification	21
4.2	FFS	21
5	Results and Analysis	23

6	Discussion	25
7	Conclusions and Future work	27
7.1	Future work	27
	References	29

List of Figures

2.1	Basic structure of inode-based filesystem	8
2.2	Simple visualization of how FUSE operations are executed . .	11
4.1	Basic structure of FFS inode-based structure	22

List of Tables

3.1	Comparison between features present in related filesystems and FFS. X means that the feature is supported and - means that it is not supported	20
-----	--	----

Listings

2.1 Pseudocode of a minimalistic inode filesystem structure	8
---	---

Chapter 1

Introduction

To keep files and data secure we often use encrypted filesystems which have been proven to be essentially unbreakable using brute-force approaches. However, while many of these filesystems hide the content of the data, they often do not conceal the existence of data. For instance, using snapshots of the filesystems from different moments in time, it could be possible to notice a difference in the data stored in these filesystems and therefore that data exists and where it is located. Snapshots could even reveal user passwords [1].

Steganographic filesystems are intended to make the data deniable, meaning that the user is supposed to be able to plausibly deny the existence of data. There are many reasons why this is important. For instance, in 2011 a Syrian man recorded videos of attacks on civilians carried out by Syrian security forces, which he wanted to share with the world [2]. By cutting his arm, he was able to hide a memory card inside the wound and smuggled it out of the country. However, if he would have used methods such as an encrypted steganographic filesystem, the border control would not have been able to discover even the existence of data, even if they would have found the memory card. By only encrypting the data, the border control could see that he was trying to hide data and force him to reveal the decryption key, either by legal measures or by force, which is why he smuggled it out.

Social media platforms such as Twitter, Flickr, and Facebook have many millions of daily users that post texts and images (for example, of their cats or funny videos). According to Henna Kermani at Twitter, they processed 200 GB of image data every second in 2016 [3]. The difference between the photos posted on Twitter compared to the ones stored on cloud services such as Google Drive is that the images on Twitter are stored for free for the user, for what seems to be an indefinite period. However, there is no obligation for

these services to save it forever, they do reserve the right to remove any content at any time, as stated in their terms of service^{*}. There is also no specified maximum lifespan of these posts. While Google Drive and similar services often have a free-tier of storage with a certain maximum storage limit, Twitter does not have a specified upper limit on how many images or tweets one can make. However, such constraints can be imposed on specific users whenever Twitter wishes as is stated in their terms of service[†].

1.1 Project Overview

This project intends to create a secure and deniable cloud-based filesystem called the *Fejk FileSystem* (FFS) which takes advantage of free online web services, such as Twitter, for the actual storage. The idea is to save the user's files by posting or sending an encrypted version of the file as posts or private messages on these web services. The intention is not to create a revolutionary fast and usable filesystem but instead to explore how well it is possible to utilize the storage that Twitter and similar services provide their users for free, as a secure and deniable cloud-based filesystem. Additionally, the performance and limits of this filesystem will be analyzed and compared to alternative filesystems, such as Google Drive, to compare the advantages and disadvantages of the developed filesystem compared to professional filesystems. The security of the filesystem will also be discussed, as well as an analysis of the steganographic capability of the developed filesystem.

1.2 Problem

Is it possible to create a secure and deniable cloud-based filesystem that stores the data on various online web services through the use of the free user accounts offered by these online web services? What are the drawbacks of such a filesystem compared to similar filesystem solutions with regards to write and read speed, storage capacity, and reliability? Are there advantages to such a system in regards to security and deniability?

^{*} <https://twitter.com/en/tos> [†] <https://twitter.com/en/tos>

1.3 Purpose and motivation

The purpose of this research is to explore the possibility to create a secure, steganographic cloud-based filesystem that stores data on online services and to compare the performance, benefits, and disadvantages of such a filesystem to existing steganographic filesystems and distributed filesystem services. A distributed filesystem service, such as Google Drive, provide data storage for users which can be both free and cost money. Even though Google Drive encrypts the user's data, they control the encryption and decryption keys, and the method of encryption [4]. This means that they can give out the user's files and data if faced with legal actions such as subpoenas. It also opens up the possibility of hackers gaining access to the files without the user having any way to control them.

The idea behind FFS is to have a cloud-based filesystem without anyone except the user being able to control the unencrypted data. We aim to exploit the storage web services give their users for free. By encrypting the files locally before uploading them to these services, the user is the only one who has access to the encryption and decryption keys and therefore unencrypted data. Even if the web service would look at the data uploaded by the user, it is not readable without the decryption key. An interesting aspect of this is that online web services, such as social media, provide users with essentially an infinite amount of storage for free. Anyone can create any number of accounts on Twitter and Facebook without cost, and with enough accounts, one could potentially store all their data using such a filesystem. As the file data is only accessible by the user, and as the filesystem can be unmounted to hide its existence, it is steganographic.

There are several steganographic filesystems available but these lack certain aspects that FFS aims to solve. Some filesystems are based on the local disk of the device in use, such as the physical storage device on a computer or phone, or an external storage device connected to a computer or phone. While such filesystems have advantages such as speed, they lack accessibility as you need to have the device to access the content on it. It also means that when you want to share or transport the data, you must physically move the device which can mean problems as it could for instance be taken from you or be destroyed. Cloud-based solutions counter this by being available from any location that has internet access to the services used. However, existing cloud-based solutions introduce other disadvantages. One example is CovertFS [5] where data is stored in images posted on web services. The images are actual images representing something meaning that there's a limit on how much

steganographic data can be stored, and the filesystems limit this to 4 kB which means that such a filesystem with a lot of data will require many images which could lead to suspicion from the owners of the web services. More examples of filesystems similar to the idea will be presented in Chapter 3.

1.4 Goals

The project aims to create a secure, deniable filesystem that stores its data on online web services by taking advantage of the storage provided to its users. This can be split into the following subgoals:

1. to create a mountable filesystem where files can be stored, read, and deleted,
2. for the filesystem to store all the data on online web services rather than on local disks,
3. for the system to be secure in the sense that even with access to the uploaded files and the software, the data is not readable without the correct decryption key,
4. to provide the user of the filesystem with reasonable deniability of its data in the sense that it is not possible to associate the user with FFS if the filesystem is not mounted,
5. to analyze the write and read speed, storage capacity, and reliability of the filesystem and compare it to commercial distributed filesystems and local filesystems, and,
6. to analyze and discuss environmental and ethical aspects of the filesystem.

1.5 Research Methodology

The filesystem created through this thesis will be developed on a Macbook laptop running macOS Monterey, version 12.0.1. It will be written in C++11 and use the Filesystem in Userspace (FUSE) MacOS library [6] which enables the writing of a filesystem in userspace rather than in kernel space. FUSE is available on other platforms too, such as Linux, but the filesystem will be developed on a Macbook laptop thus macFUSE is chosen. C++ is chosen

because the FUSE API is available in C, and C++ version 11 is well established and used. Further details about the development environment is found in Section 4.1.

The resulting filesystem will be evaluated against other filesystems, both commercial distributed systems, such as Google drive, and an instance of Apple File System (APFS) [7] on the Macbook laptop referenced above. Quantitative data will be gathered from the different filesystems through the use of experiments with the filesystem benchmarking software IOzone [8]. IOzone was chosen because it is, compared to tools such as Fio and Bonnie++, simpler to use while still powerful [9]. We will look at attributes such as the differences in read and write speeds between different filesystems, as well as the speed of random read and random write. However, according to Tarasov, Bhanage, *et al.*, benchmarking filesystems using benchmarking tools is difficult to perform in a standardized way [10] which will be taken into consideration during the evaluation and when concluding the thesis.

1.6 Delimitations

Due to limitations in time and as the system is only a prototype for a working filesystem and not a production filesystem, some features found in other filesystems are not going to be implemented in FFS. The focus will be to implement a subset of the POSIX standard functions, containing only crucial functions for a simple filesystem, specifically, the FUSE functions *open*, *read*, *write*, *mkdir*, *rmdir*, *readdir*, and *rename*. However, file access control is not a necessity and will therefore not be implemented, thus functions such as *chown* and *chmod* are not going to be implemented. The reason is that the goal is to present and evaluate the possibility of creating a secure steganographic filesystem with a storage medium based on online web services and thus FFS will only aim to implement a minimal filesystem.

There is also an argument that could be made that FFS should support multiple users so that anyone can mount the filesystem but only browse their files. However, as this project is only a proof-of-concept of the filesystem, this will not be implemented. Instead, FFS will be built for single-user support where only a password will unlock everything FFS is storing. This means that anyone who mounts FFS with the password will access everything that other users might have stored.

1.7 Structure of the thesis

Chapter 2 presents theoretical background information of filesystems and the basis of FFS while Chapter 3 mentions and analyses related work. Chapter 4 describes the implementation and the design choices made for the system, along with the analysis methodology. Chapter 5 presents the results of the analysis and Chapter 6 discusses the findings and other aspects of the work. Lastly, Chapter 7 will finalize the conclusion of the thesis and discuss potential future work.

Chapter 2

Background

This chapter presents concepts and information that is relevant for understanding, implementing, and evaluating FFS. We first present the idea of inode-based filesystems and how data is stored in a filesystem. Following is the introduction of Filesystem in Userspace (FUSE) which will be used to implement the filesystem. Later sections present background information about Twitter and the potential threat adversaries of the filesystem.

2.1 Filesystems and data storage

This section presents how certain filesystems used today are structured. We present the idea of inode-based filesystems and distributed filesystems. Following, we describe how data is stored in a storage system and how this information can be used in FFS.

2.1.1 Unix filesystems

A Unix filesystem uses a data structure called an *inode*. The inodes are found in an inode table and each inode keeps track of the size, blocks used for the file's data, and metadata for the files in the filesystem. A directory simply contains the file names and each file or directory's inode id. The system can with an inode id find information about the file or directory using the inode table. Each inode can contain any metadata that might be relevant for the system, such as creation time and last update time.

Figure 2.1 shows how example inode filesystem and how it can be visualized. The blocks of an inode entry are where in the storage device the data is stored, each block is often defined as a certain amount of bytes.

Listing 2.1 describes a simple implementation of an inode, an inode table, and directory entries.




Figure 2.1: Basic structure of inode-based filesystem

Listing 2.1: Pseudocode of a minimalistic inode filesystem structure

```
struct inode_entry {
    int      length
    int []   blocks
    // Metadata attributes are defined here
}

struct directory_entry {
    char*    filename
    int      inode
}

// Maps inode_id to a inode_entry
map<int, inode_entry> inode_table
```

Different filesystems provide different features and limitations. The Extended Filesystem (ext) exists in four different versions: ext, ext2, ext3, and ext4. This filesystem is often used on Unix systems. Each iteration brings new

features and changes the limitations. For instance, comparing the two latest iterations, ext3 and ext4. This filesystem can theoretically store files up to 16 TiB while ext3 can store files up to 2 TiB [11]. Additionally, ext4 supports timestamps in units of nanoseconds while ext3 only supports timestamps with a resolution of one second. ext4 natively supports encryption at the directory level through the use of the fscrypt API [12].

The Apple Filesystem (APFS) is a modern filesystem that is used on iPhones and Mac and can store files with a size up to 9 EB [13]. It supports timestamps in units of nanoseconds and is built to be used on solid-state drives (SSD) [14]. It also supports modern features that its predecessor Mac OS Extended (HFS+) does not support, such as Snapshots and Space Sharing. APFS natively supports encryption of the filesystem volume [15].

2.1.2 Distributed filesystems

Filesystems are used to store data on for instance a hard drive of a computer locally or in the cloud. For example, Google Drive is a filesystem that enables users to save their data online with up to 15 GB for free [16] using Google's clusters of distributed storage devices, meaning that the data is saved on Google's servers which can be located wherever they have data centers [17]. Paying customers can have a greater amount of storage using the service. Apple's iCloud and Microsoft's OneDrive are two additional examples of distributed filesystems where users have the option of free-tier and paid-tier storage.

Cloud-based filesystems, as opposed to a filesystem on a physical disk, are accessible from multiple computers and devices without requiring the user to connect a physical disk to the computer. Instead, as the filesystem is accessible through the internet, it can be accessed regardless of the user's location or device as long as a connection to the filesystem can be established. Thus, even if the user would lose their computer or if it would malfunction, the data on the cloud-based filesystem can still be accessed which means that the data could still be recovered.

2.1.3 Data storage and encoding

Different file types have different protocols and definitions of how they should be encoded and decoded, for instance, a JPEG and a PNG file can be used to display similar content but the data they store is different. At the lowest level, storage devices often represent files as a string of binary digits no matter

the file type (however there are non-binary storage devices [18], but this is outside the scope of this thesis). If one would represent an arbitrary file of X bytes, each byte (0x00 - 0xFF) can be represented as a character such as the Extended ASCII (EASCII) keyset and we can therefore decode this file as X different characters. Using the same set of characters for encoding and decoding we can get a symmetric relation for representing a file as a string of characters. EASCII is only one example of such a set of characters, any set of strings with 256 unique symbols can be used to create such a symmetric relation, for instance, 256 different emojis or a list of 256 different words. However, if we are using a set of words we could also have to introduce a unique separator so that the words can be distinguished. If we would use a single space character as the separator, we could make the encoded text look like a text document; however, random words one after another lead to a high probability of creating an unstructured text document. Further, if punctuation is introduced, for instance as part of some words, the text document could look like it contains random and unstructured sentences.

This string of X bytes can also be used as the data in an image. An image can be abstracted as a $h * w$ matrix, where each element is a pixel of a certain color. In an image with 8-bit Red-Green-Blue (RGB) color depth, each pixel consists of three 8-bit values, i.e. three bytes. One can therefore imagine that we can use this string of X bytes to assign colors in this pixel matrix by assigning the first three bytes as the first pixel's color, the next three bytes as the following pixel's color, and so forth. This means that X bytes of data can be represented as

$$\text{ceil}(\frac{X}{3})$$

pixels, where *ceil* rounds a float to the closest larger integer. For a file of 1 MB, i.e. $X = 1\,000\,000$ we need 333 334 pixels in an image with 8-bit RGB color depth. The values of h and w are arbitrary but if we for instance want a square image we can set $h = w = 578$ which means that there will be 334 084 pixels in total, and the remaining 750 pixels will just be fillers to make the image a reasonable size. Using filler pixels requires us to keep track of the number of bytes that we store in the image so that we do not read the filler bytes when the image is decoded. However, we could choose $h = 1$ and $w = 333\,334$ which would mean a very wide image but would not require filler pixels.

This means that we can represent any file as a string of bytes which can then be encoded into text or as an image, which can be posted on for instance social media. However, there is a possibility that the social media services compress the images uploaded which could lead to data loss in the image, which would

mean that the decoded data would be different from the encoded data. In this case, we would not be able to retrieve the original data that was stored.

2.2 FUSE

Filesystem in Userspace (FUSE) is a library that provides an interface to create filesystems in userspace rather than in kernel space which is otherwise often considered the standard when writing commercial filesystems [19]. The reason to implement a filesystem in kernel space is that it leads to faster system calls than when writing a filesystem in userspace. However, while filesystems written with FUSE are generally slower than a kernel-based filesystem, using FUSE simplifies the process of creating filesystems. macFUSE is a port of FUSE that operates on Apple's macOS operating system and it extends the FUSE API [6]. macFUSE provides an API for C and Objective C.

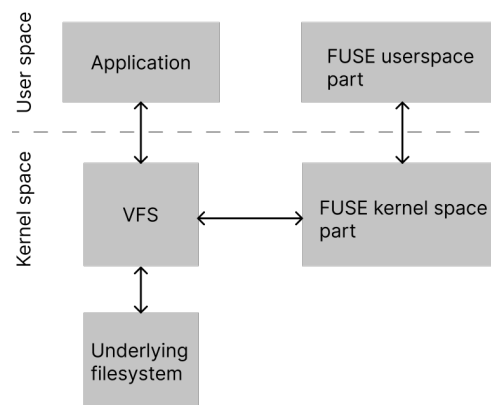


Figure 2.2: Simple visualization of how FUSE operations are executed

Figure 2.2 shows an overview how FUSE works. FUSE consists of a kernel space part and a userspace part that perform different tasks [20]. The kernel part of FUSE operates with the Virtual Filesystem (VFS) which is a layer in both the Linux kernel and the macOS kernel that exposes a filesystem interface for userspace applications [21, 22]. The VFS interface is independent of the underlying filesystem and is an abstraction of the underlying filesystem operations which can be used on any filesystem the VFS supports. The userspace part of FUSE communicates with the kernel space part through a block device. Operations on a mounted FUSE filesystem are sent to the VFS from the user application, which is then sent to the kernel part of FUSE. If needed, the operations are transmitted to the userspace part of FUSE where

the operation is handled and a response is sent back to the VFS and the user application through the FUSE kernel module. However, some actions can be handled by the FUSE kernel module directly, such as if the file is cached in the kernel part of FUSE [20]. The response is then sent back to the user application from the kernel module through the VFS.

2.3 Twitter

Twitter is a micro-blog online where users can sign up for a free account and create public posts (tweets) using text, images, and videos. Text posts are limited to 280 characters while images can be up to 5 MB and videos up to 512 MB [23]. There is also a possibility to send private messages to other accounts, where each message can contain up to 10 000 characters and the same limitations on files. However, direct messages older than 30 days are not possible to retrieve through Twitter's API [24]. It is possible to create threads of Twitter posts where multiple tweets can be associated in chronological order.

Twitter's API defines technical limits of how many times certain actions can be executed by a user [25]. A maximum of 2 400 tweets can be sent per day, and the limit is further broken down into smaller limits at semi-hourly intervals. Hitting a limit means that the user account no longer can perform the actions that the limit represents until the time period has elapsed.

2.4 Threats

To consider a filesystem secure it is important to imagine different potential adversaries who might attack the system. Considering that FFS has no real control of the data stored on the different services, all the data must be considered to be stored in an insecure system. Even if we could hide the posts made on for instance Twitter by making the profile private, we must still consider that Twitter themselves could be an adversary or that they could potentially give out information, such as tweets or direct messages, to entities such as the police. Twitter's privacy policy mentions that they may share, disclose, and preserve personal information and content posted on the service, even after account deletion for up to 18 months [26]. Therefore, to achieve security the data stored must always be encrypted. We assume that an adversary has access to all knowledge about FFS, including how the data is converted, encrypted, and posted. We also assume they know which websites

and accounts could post data from the filesystem - but we assume they do **not** have the decryption key. There are multiple secure ways of encrypting data, including AES which is one of the faster and more secure encryption algorithms [27]. However, even though the data is encrypted, other properties such as your IP address can be compromised which can expose the user's identity. The problem of these other sources of information external to FFS is not addressed in FFS but remains for future work.

Other than adversaries for FFS, we might also imagine that the underlying services might face attacks that can potentially harm the security of the system or even cause the service to go offline, potentially indefinitely. One solution is to use redundancy - by duplicating the data over multiple services, we can more confidently believe that our data will be accessible as the probability of all services going offline at the same time is lower.

The deniability of FFS is an important aspect of the filesystem. Potential threat adversaries are agents that the user is trying to hide the data from, such as governing states. For the system to be deniable, an adversary should not be able to gain any information about anything about the potential data in the system, this includes even the existence of data. When the filesystem is unmounted there should be no trace of the filesystem ever being present in the device. We will assume that an adversary is competent and can analyze the software and hardware completely.

Chapter 3

Related work

The research area of creating filesystems to improve security, reliability, and deniability is not new and has been well worked on previously. This chapter presents previous work that is related to this thesis. This includes other filesystems that share similarities with the idea of FFS, for instance within the idea of unconventional storage media and the area of steganography.

3.1 Steganography and deniable filesystems

Steganography is the art of hiding information in plain sight and has been around for ages. Today, a major part of steganography is hiding malicious code in for instance images, called stegomalware or stegoware. Stegomalware is an increasing problem and in a sample set of examined real-life stegomalware, over 40% of the cases used images to store the malicious code [28]. While FFS will not include malicious code in its images, this stegomalware problem has fostered the development of detection techniques of steganography in for instance social media, and it is well researched.

Twitter has been exposed to allowing steganographic images that contain any type of file easily [29]. David Buchanan created a simple python script of only 100 lines of code that can encode zip-files, mp3-files, and any file imaginable in an image of the user's choosing [30]. He presents multiple examples of this technique on his Twitter profile*. The fact that the images are available for the public's eye might be evidence that Twitter's steganography detection software is not perfect. However, it is also possible that Twitter has chosen to not remove these posts.

* <https://twitter.com/David3141593>

Other examples of steganographic data storage on Open Social Networks (OSNs) include the paper presented by Ning, Singh, *et al.* where the authors build a system for private communication on public photo-sharing web services [31]. Due to the web services processing of uploaded multimedia, they first researched how the integrity of steganographic data could be maintained after being uploaded to these services. Following this, they presented an approach that ensured the integrity of the hidden messages in the uploaded images, while also maintaining a low likelihood of discovery from the steganographic analysis. Beato, De Cristofaro, and Rasmussen also explores the idea of undetectable communications over OSNs in another paper [32]. While implementation is not carried out, they present an idea where messages are encoded together with a cover object and a cryptographic key to produce a steganographic message which is then posted to the OSN. A web-based user interface client with a PHP server backend is presented as the method the users would use to create and share their secret messages.

A steganographic, or deniable, filesystem is a system that does not expose files stored on this system without credentials - neither how many files are stored, their sizes, their content, or even if there exist any files in the filesystem [33]. This is also known as a rubber hose filesystem because of the characteristic that the data only can be proven to exist with the correct encryption key which only is accessible if the person is tortured and beaten with a rubber hose because of its simplicity and immediacy compared to the complexity of breaking the key by computational techniques.

3.2 Cryptography

Some papers choose to invent their encryption methods rather than using established standards. Chuman, Sirichotedumrong, and Kiya proposes a scrambling-based encryption scheme for images that splits the picture into multiple rectangular blocks that are randomly rotated and inverted, both horizontally and vertically, along with shuffling of the color components [34]. This is used to demonstrate the security and integrity of images sent over insecure channels. The paper uses Twitter and Facebook to exhibit this. Despite its improvement and compatibility of a common image format, such as bitstream compliance, due to its well-proven security FFS will use AES as its encryption method.

3.3 Related filesystems

Multiple steganographic filesystems have been presented previously but many of these are focused on filesystems for physical storage disks that the user has access to. For instance, Timothy Peters created DEFY, a deniable filesystem using a log-based structure in 2014 [35]. DEFY was built to be used exclusively on Solid State Drives (SSD) found in mobile devices to provide a steganographic filesystem that could be used on Android phones. Further examples of local disk-based filesystems can be found in [1, 33, 36, 37], among other papers. However, this paper aims to create a filesystem that is not based on a physical disk but rather a cloud-based steganographic filesystem that uses online web services as its storage medium.

In 2007, Baliga, Kilian, and Iftode presented an idea of a covert filesystem that hides the file data in images and uploads them to web services, named CovertFS [5]. The paper lacks implementation of the filesystem but they present an implementation plan which includes using FUSE. They limit the filesystem such that each image posted will only store a maximum of 4 kB of steganographic file data and the images posted on the web services will be actual images. This is different from the idea of FFS where the images will be purely the encrypted file data and will therefore not be an image that represents anything but will instead look like random color noise. An implementation of CovertFS has been attempted by Sosa, Sutton, and Huang which also used Tor to further anonymize the users [38].

In 2016, Szczypiorski introduced the idea of StegHash - a way to hide steganographic data on Open Social Networks (OSN) by connecting multimedia files, such as images and videos, with hashtags [39]. Specifically, images were posted to Twitter and Instagram along with certain permutations of hashtags that pointed to other posts through the use of a custom-designed secret transition generator. StegHash managed to store short messages with 10 bytes of hidden data with a 100% success rate, while longer messages with up to 400 bytes of hidden data had a success rate of 80%. Bieniasz and Szczypiorski later presented SocialStegDisc which was a filesystem application of the idea presented with StegHash [40]. Multiple posts could be required to store a single file and each post referenced the next post like a linked list, which means that you only need the root post to read all the data. This is unlike the idea of FFS where a table will be kept to keep track of which posts store a certain file, and in what order they should be concatenated. SocialStegDisc lacks actual implementation of the filesystem but similar to CovertFS presents the idea of a social media-based filesystem.

TweetFS is a filesystem created by Robert Winslow that stores the data on Twitter [41], created in 2011. It was created as a proof of concept to show that it is possible to store file data on Twitter. The filesystem uses sequential text posts to store the data. The filesystem is not mounted to the operating system, instead, the user interacts with a Python script through the command line. This makes the filesystem less convenient from a user perspective, compared to a mounted filesystem where the files can be browsed using a user interface or command line. There are two commands available: `upload` and `download` which upload and download files or directories, respectively. Names and permissions of files and directories are maintained throughout the upload and download process. The tweets are not encrypted but are enciphered into English words which makes them look like nonsense paragraphs, similar to what we mentioned in Section 2.1.3 about how arbitrary data can be encoded as plain text. This makes the filesystem less secure than an encrypted version as it can be read by anyone with access to the decoder. However, it does introduce a steganographic element to the filesystem.

In 2006, Jones created GmailFS - a mountable filesystem that uses Google's Gmail to store the data [42, 43]. The filesystem was written in Python using FUSE and was presented well before the introduction of Google Drive in 2012. It does not support encryption as the plain file data is stored in emails. Today, Gmail and Google Drive share their storage quota and GmailFS has since become redundant as Google drive is an easier filesystem to use. GMail Drive is another example of a Gmail-based filesystem and it was influenced by GmailFS [44]. GMail Drive has been declared dead by its author since 2015.

Badulescu, Shender, and Zadok created Cryptfs, a stackable Vnode filesystem that encrypted the underlying, potentially unencrypted, filesystem [45]. By making the filesystem stackable, any layer can be added on top of any other, and the abstraction occurs by each Vnode layer communicating with the one beneath. There is a potential to further stack additional layers by using tools such as FiST [46]. This approach enables one to create not only an encrypted file system but also to provide redundancy by replicating data to different underlying filesystems. If these filesystems are independent, then potentially this increases availability and reliability. FFS aims to achieve stackability through the use of FUSE.

3.4 Filesystem benchmarking

IOzone is a filesystem benchmarking tool that is used to measure performance and analyze a filesystem [8]. It is built for, among other platforms, Apple's

macOS where the filesystem will be built, run, and tested. However, as mentioned previously, filesystem benchmarking is more complicated than one might imagine. Different filesystems might perform differently on small and big file sizes among other things, which means that we can never compare benchmarking outputs as just single numbers. We must instead compare different aspects of the filesystems. In 2011 Tarasov, Bhanage, *et al.* presents a paper where they criticize several papers due to their lack of scientific and honest filesystem benchmarking [10]. The problem of benchmarking a filesystem is all the different components that are involved when interacting with a filesystem. For instance, they mention how benchmarking the in- and output (I/O) of the filesystem, such as bandwidth and latency, is different from benchmarking on-disk operations, such as the performance of file read and write operations. The benchmarking tools can rarely affect or determine how the filesystem handles caching and pre-fetching. This means that benchmarking the read and write performance of different filesystems can be misleading as they might handle this differently, meaning that the result could be different depending on for instance the distance between the files on the disk. Two files could be adjacent on the disk on one filesystem and therefore one could be pre-fetched into the cache when the other one is read. Considerations about such factors must be present when analyzing the results of the benchmarking.

Tarasov, Bhanage, *et al.* also lists several different filesystem benchmarking tools available and used by the papers they reviewed, and how well the tools can analyze certain aspects of a filesystem [10]. IOZone is listed as being compatible with multiple of the different benchmarking types and as it is simpler to use [9] and still maintained, this was chosen as the benchmarking tool for FFS.

3.5 Summary

As presented, different filesystems provide different features and drawbacks. In Table 3.1 we display a summary of characteristics and features of some filesystems mentioned above and how FFS compares. As can be seen, FFS mainly lacks certain filesystem operations which are not the focus of FFS as it is a proof of concept.

	ext4	Google drive	DEFY	TweetFS	FFS
Mountable	X	X	X	-	X
Read/Write/Remove file	X	X	X	X	X
Read/Write/Remove directory	X	X	X	X	X
Hard links	X	-	X	-	-
Soft links	X	-	X	-	-
File and directory access control	X	X	-	X	-
Encrypted	X	X*	X	-	X
Steganographic	-	-	X	X	X
Cloud-based	-	X	-	X	X

*As mentioned, the user has no control over this encryption

Table 3.1: Comparison between features present in related filesystems and FFS. X means that the feature is supported and - means that it is not supported

Chapter 4

Method

This section presents the methodology of implementing FFS and the specifications of the development environment. We also present how the quantitative data used for the evaluation is acquired. Also, the experiments on the filesystems are presented.

4.1 Development environment specification

4.2 FFS

The artifact that will be developed as a result of this thesis is the Fejk FileSystem (FFS). It uses online services to store the data but behaves as a mountable filesystem for the users. The filesystem will be minimal and not support all functionalities that other filesystems do, such as links. The reasoning is that these behaviors are not required for a useable system, and when comparing the system to distributed filesystems such as Google Drive, many of these other filesystems also often do not support links.

Figure 4.1 presents the basic outline of FFS and a example content of the filesystem. FFS is based on the idea of inode filesystems but instead of an inode pointing to specific blocks in a disk, the inodes of FFS will instead keep track of the id numbers of the posts on the online services where the file is located.

As a deniable filesystem, you should not be able to gain any information about the filesystem without the correct credentials. Using incorrect credentials, the filesystem will just show an empty filesystem meaning that the adversary will know if the credentials provided are correct and that the

filesystem is empty, or if the credentials were incorrect. Only when the correct credentials are provided will FFS show the filesystem of the user.

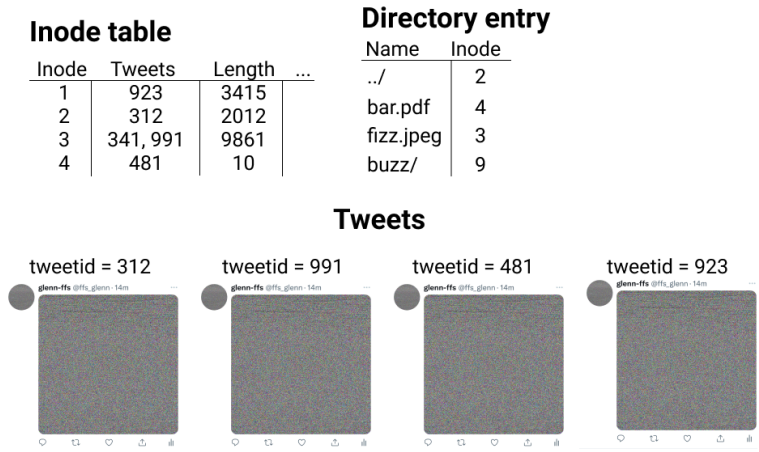


Figure 4.1: Basic structure of FFS inode-based structure

Chapter 5

Results and Analysis

Chapter 6

Discussion

Chapter 7

Conclusions and Future work

This chapter presents the conclusions from the thesis from what has been discussed under Chapter 6. Finally, future work on the topic is discussed.

7.1 Future work

As mentioned previously, FFS does not implement all features that the POSIX standard defines. Future development for FFS could be to implement more of these functions, such as links and file permissions. This could make the filesystem resemble a regular filesystem further. Another improvement could be to move from userspace using FUSE, to kernel space. This could speed up filesystem operations. Another feature that could be interesting to evaluate is the possibility to share files with other users, similar to Google Drive.

Even though the files are encrypted so that the data is confidential, further research could include hiding the user's online activity through the use of for instance Tor. Currently, the integrity of the user is not considered but for the filesystem to be further plausibly deniable, this should be addressed as the user could otherwise be identified by its IP address and other online fingerprints that could be provided by the online web services.

To improve the dependability of the filesystem, support for more online web services could be implemented. For instance, Github provides free user accounts with many gigabytes of data. Even free-tier distributed filesystems, such as Google Drive, could be utilized. If multiple user accounts are used in coordination over multiple services, the filesystem could achieve even more storage.

References

- [1] J. Han, M. Pan, *et al.*, “A multi-user steganographic file system on untrusted shared storage”, in *Proceedings of the 26th Annual Computer Security Applications Conference on - ACSAC '10*, Austin, Texas: ACM Press, Dec. 6, 2010, p. 317, ISBN: 978-1-4503-0133-6. DOI: [10.1145/1920261.1920309](https://doi.org/10.1145/1920261.1920309). [Online]. Available: <http://portal.acm.org/citation.cfm?doid=1920261.1920309> (visited on 01/27/2022).
- [2] R. Westhead, “How a Syrian refugee risked his life to bear witness to atrocities”, *The Toronto StarWorld*, Mar. 14, 2012, ISSN: 0319-0781. [Online]. Available: https://www.thestar.com/news/world/2012/03/14/how_a_syrian_refugee_risked_his_life_to_bear_witness_to_atrocities.html (visited on 04/13/2022).
- [3] “Mobile @Scale London recap - Engineering at Meta”. (), [Online]. Available: <https://engineering.fb.com/2016/03/29/android/mobile-scale-london-recap/>.
- [4] D. Johnson. “Is Google Drive secure? How Google uses encryption to protect your files and documents, and the risks that remain”, *Business Insider*. (Feb. 25, 2021), [Online]. Available: <https://www.businessinsider.com/is-google-drive-secure> (visited on 04/13/2022).
- [5] A. Baliga, J. Kilian, and L. Iftode, “A web based covert file system”, in *Proceedings of the 11th USENIX Workshop on Hot Topics in Operating Systems*, ser. HOTOS'07, USA: USENIX Association, May 7, 2007.
- [6] “Home - macFUSE”. (), [Online]. Available: <https://osxfuse.github.io/> (visited on 03/07/2022).

- [7] Apple. “About Apple File System | Apple Developer Documentation”. (), [Online]. Available: https://developer.apple.com/documentation/foundation/file_system/about_apple_file_system (visited on 03/13/2022).
- [8] “Iozone Filesystem Benchmark”. (), [Online]. Available: <https://www.iozone.org/> (visited on 03/07/2022).
- [9] U. K. Agarwal. “Comparing IO benchmarks: FIO, IOZONE and BONNIE++”, FuzzyWare. (May 19, 2018), [Online]. Available: <https://uditagarwal.in/comparing-io-benchmarks-fio-iozone-and-bonnie/> (visited on 03/13/2022).
- [10] V. Tarasov, S. Bhanage, *et al.*, “Benchmarking file system benchmarking: It *IS* rocket science”, in *13th Workshop on Hot Topics in Operating Systems (HotOS XIII)*, Napa, CA: USENIX Association, May 2011. [Online]. Available: <https://www.usenix.org/conference/hotosxiii/benchmarking-file-system-benchmarking-it-rocket-science>.
- [11] J. Salter. “Understanding Linux filesystems: ext4 and beyond”, Opensource.com. (Apr. 2, 2018), [Online]. Available: <https://opensource.com/article/18/4/ext4-filesystem> (visited on 03/09/2022).
- [12] “Fscrypt - ArchWiki”. (), [Online]. Available: <https://wiki.archlinux.org/title/Fscrypt> (visited on 04/25/2022).
- [13] “APFS (Apple File System) Key Features | iGotOffer”, About Apple | iGotOffer. (Jul. 16, 2017), [Online]. Available: <https://igotoffer.com/apple/apfs-apple-file-system-key-features> (visited on 04/11/2022).
- [14] T. Nelson. “What Is APFS and Does My Mac Support the New File System?”, Lifewire. (), [Online]. Available: <https://www.lifewire.com/apple-apfs-file-system-4117093> (visited on 04/11/2022).
- [15] A. Inc. “File system formats available in Disk Utility on Mac”, Apple Support. (), [Online]. Available: <https://support.apple.com/guide/disk-utility/file-system-formats-dskul9ed921c/mac> (visited on 04/25/2022).

- [16] “Cloud Storage for Work and Home – Google Drive”. (), [Online]. Available: <https://www.google.com/intl/sv/drive/> (visited on 10/26/2021).
- [17] “Distributed Storage: What’s Inside Amazon S3?”, Cloudian. (), [Online]. Available: <https://cloudian.com/guides/data-backup/distributed-storage/> (visited on 10/26/2021).
- [18] “Multi-state data storage leaving binary behind: Stepping ‘beyond binary’ to store data in more than just 0s and 1s”, ScienceDaily. (Oct. 12, 2020), [Online]. Available: <https://www.sciencedaily.com/releases/2020/10/201012115937.htm> (visited on 03/10/2022).
- [19] *Libfuse*, libfuse, Oct. 26, 2021. [Online]. Available: <https://github.com/libfuse/libfuse> (visited on 10/26/2021).
- [20] B. K. R. Vangoor, V. Tarasov, and E. Zadok, “To {FUSE} or Not to {FUSE}: Performance of {User-Space} File Systems”, presented at the 15th USENIX Conference on File and Storage Technologies (FAST 17), Feb. 27–Mar. 2, 2017, pp. 59–72, ISBN: 978-1-931971-36-2. [Online]. Available: <https://www.usenix.org/conference/fast17/technical-sessions/presentation/vangoor> (visited on 04/06/2022).
- [21] R. Gooch. “Overview of the Linux Virtual File System — The Linux Kernel documentation”. (), [Online]. Available: <https://www.kernel.org/doc/html/latest/filesystems/vfs.html> (visited on 04/12/2022).
- [22] A. Singh, *Mac OS X Internals: A Systems Approach*. Pearson, 2006, ISBN: 0-321-27854-2. [Online]. Available: <https://flylib.com/books/en/3.126.1.136/1/> (visited on 04/11/2022).
- [23] “Media Best Practices - Twitter”. (), [Online]. Available: <https://developer.twitter.com/en/docs/twitter-api/v1/media/upload-media/uploading-media/media-best-practices> (visited on 10/26/2021).
- [24] “Retrieving older than 30 days Direct Messages (direct_messages/events/list) - Twitter API / Standard APIs v1.1”, Twitter Developers. (Apr. 27, 2018), [Online]. Available: <https://twittercommunity.com/t/retrieving-older-than-30-days-direct-messages-direct-messages-events-list/104901> (visited on 03/11/2022).

- [25] “Understanding Twitter limits | Twitter Help”. (), [Online]. Available: <https://help.twitter.com/en/rules-and-policies/twitter-limits> (visited on 03/11/2022).
- [26] Twitter. “Privacy Policy”. (), [Online]. Available: <https://twitter.com/en/privacy> (visited on 02/15/2022).
- [27] P. Mahajan and A. Sachdeva, “A Study of Encryption Algorithms AES, DES and RSA for Security”, *Global Journal of Computer Science and Technology*, Dec. 7, 2013, ISSN: 0975-4172. [Online]. Available: <https://computerresearch.org/index.php/computer/article/view/272> (visited on 02/07/2022).
- [28] S. C. Foundation. “SIMARGL: Stegware primer, part 1”. (Feb. 14, 2020), [Online]. Available: <https://cuing.eu/blog/technical/simargl-stegware-primer-part-1> (visited on 02/09/2022).
- [29] “Twitter images can be abused to hide ZIP, MP3 files — here’s how”. (), [Online]. Available: <https://www.bleepingcomputer.com/news/security/twitter-images-can-be-abused-to-hide-zip-mp3-files-heres-how/> (visited on 02/09/2022).
- [30] D. Buchanan, *Tweetable-polyglot-png*, Feb. 9, 2022. [Online]. Available: <https://github.com/DavidBuchanan314/tweetable-polyglot-png> (visited on 02/09/2022).
- [31] J. Ning, I. Singh, *et al.*, “Secret message sharing using online social media”, in *2014 IEEE Conference on Communications and Network Security*, Oct. 2014, pp. 319–327. doi: [10.1109/CNS.2014.6997500](https://doi.org/10.1109/CNS.2014.6997500).
- [32] F. Beato, E. De Cristofaro, and K. B. Rasmussen, “Undetectable communication: The Online Social Networks case”, in *2014 Twelfth Annual International Conference on Privacy, Security and Trust*, Jul. 2014, pp. 19–26. doi: [10.1109/PST.2014.6890919](https://doi.org/10.1109/PST.2014.6890919).
- [33] R. Anderson, R. Needham, and A. Shamir, “The Steganographic File System”, in *Information Hiding*, D. Aucsmith, Ed., ser. Lecture Notes in Computer Science, Berlin, Heidelberg: Springer, 1998, pp. 73–82, ISBN: 978-3-540-49380-8. doi: [10.1007/3-540-49380-8_6](https://doi.org/10.1007/3-540-49380-8_6).

- [34] T. Chuman, W. Sirichotedumrong, and H. Kiya, “Encryption-Then-Compression Systems Using Grayscale-Based Image Encryption for JPEG Images”, *IEEE Transactions on Information Forensics and Security*, vol. 14, no. 6, pp. 1515–1525, Jun. 2019, issn: 1556-6021. DOI: [10.1109/TIFS.2018.2881677](https://doi.org/10.1109/TIFS.2018.2881677).
- [35] T. M. Peters, “DEFY: A Deniable File System for Flash Memory”, California Polytechnic State University, San Luis Obispo, California, Jun. 1, 2014. DOI: [10.15368/theses.2014.76](https://doi.org/10.15368/theses.2014.76). [Online]. Available: <http://digitalcommons.calpoly.edu/theses/1230> (visited on 10/19/2021).
- [36] A. D. McDonald and M. G. Kuhn, “StegFS: A Steganographic File System for Linux”, in *Information Hiding*, A. Pfitzmann, Ed., ser. Lecture Notes in Computer Science, Berlin, Heidelberg: Springer, 2000, pp. 463–477, ISBN: 978-3-540-46514-0. DOI: [10.1007/10719724_32](https://doi.org/10.1007/10719724_32).
- [37] J. Domingo-Ferrer and M. Bras-Amorós, “A Shared Steganographic File System with Error Correction”, in *Modeling Decisions for Artificial Intelligence*, V. Torra and Y. Narukawa, Eds., ser. Lecture Notes in Computer Science, Berlin, Heidelberg: Springer, 2008, pp. 227–238, ISBN: 978-3-540-88269-5. DOI: [10.1007/978-3-540-88269-5_21](https://doi.org/10.1007/978-3-540-88269-5_21).
- [38] C. Sosa, B. Sutton, and H. Huang, “The Super Secret File System”, 2007. [Online]. Available: <https://www.cs.virginia.edu/~evans/wass/projects/ssfs.pdf> (visited on 03/09/2022).
- [39] K. Szczypiorski, “StegHash: New Method for Information Hiding in Open Social Networks”, *International Journal of Electronics and Telecommunications*; 2016; vol. 62; No 4, 2016, issn: 2300-1933. [Online]. Available: <https://journals.pan.pl/dlibra/publication/116930/edition/101655> (visited on 04/13/2022).
- [40] J. Bieniasz and K. Szczypiorski, “SocialStegDisc: Application of steganography in social networks to create a file system”, in *2017 3rd International Conference on Frontiers of Signal Processing (ICFSP)*, Sep. 2017, pp. 76–80. DOI: [10.1109/ICFSP.2017.8097145](https://doi.org/10.1109/ICFSP.2017.8097145).
- [41] R. Winslow. “Tweetfs/tweetfs at master · rw/tweetfs”, GitHub. (). [Online]. Available: <https://github.com/rw/tweetfs> (visited on 04/06/2022).

- [42] R. Jones. “Google hack: Use Gmail as a Linux filesystem”, Computerworld. (Sep. 15, 2006), [Online]. Available: <https://www.computerworld.com/article/2547891/google-hack--use-gmail-as-a-linux-filesystem.html> (visited on 03/09/2022).
- [43] R. Jones. “Gmail Filesystem Implementation Overview”. (Apr. 11, 2006), [Online]. Available: <https://web.archive.org/web/20060411085901/http://richard.jones.name/google-hacks/gmail-filesystem/gmail-filesystem-implementation.html> (visited on 03/09/2022).
- [44] B. Viksoe. “Viksoe.dk - GMail Drive shell extension”. (Apr. 10, 2004), [Online]. Available: <http://www.viksoe.dk/code/gmail.htm> (visited on 03/09/2022).
- [45] I. Badulescu, A. Shender, and E. Zadok, “Cryptfs: A Stackable Vnode Level Encryption File System”, 1998. doi: 10.7916/D82N5935. [Online]. Available: <https://doi.org/10.7916/D82N5935> (visited on 03/04/2022).
- [46] “FiST: Stackable File System Language and Templates”. (), [Online]. Available: <https://www.filesystems.org/> (visited on 02/02/2022).

For DIVA

```
{
  "Author1": { "Last name": "Olsson",
    "First name": "Glenn",
    "Local User Id": "u18orpa8",
    "E-mail": "glennol@kth.se",
    "organisation": {"L1": "School of Electrical Engineering and Computer Science",
    }
  },
  "Degree1": {"Educational program": "Master's Programme, Computer Science, 120 credits"
    , "programcode": "TCSCM"
    , "Degree": "Degree of Master (120 credits)"
    , "subjectArea": "Computer Science and Engineering"
  },
  "Title": {
    "Main title": "FFS: A secure cloud-based steganographic filesystem through exploitation of online web services",
    "Subtitle": "Store your sensitive data in plain sight",
    "Language": "eng" },
    "Alternative title": {
      "Main title": "FFS: Ett säkert molnbaserat steganografiskt filsystem genom utnyttjande av onlinebaserade webbtjänster",
      "Subtitle": "Lagra din känsliga data öppet",
      "Language": "swe"
    },
    "Supervisor1": { "Last name": "Ghasemirahni",
      "First name": "Hamid",
      "Local User Id": "u1fz5jtv",
      "E-mail": "hamidgr@kth.se",
      "organisation": {"L1": "",
        "L2": "Computer Science" }
    },
    "Supervisor2": { "Last name": "Peterson",
      "First name": "Zachory",
      "E-mail": "znjpeter@kth.se",
      "Other organisation": "Cal Poly"
    },
    "Examiner1": { "Last name": "Maguire Jr",
      "First name": "Gerald Quentin",
      "Local User Id": "u1d13i2c",
      "E-mail": "maguire@kth.se",
      "organisation": {"L1": "",
        "L2": "Computer Science" }
    },
    "National Subject Categories": "10201",
    "Other information": {"Year": "2022", "Number of pages": "xv,35"},
    "Series": { "Title of series": "TRITA-EECS-EX", "No. in series": "2022:00" },
    "Opponents": { "Name": "A. B. Normal & A. X. E. Normalé"},
    "Presentation": { "Date": "2022-03-15 13:00"
      , "Language": "eng"
      , "Room": "via Zoom https://kth-se.zoom.us/j/ddddeeeeee"
      , "Address": "Isafjordsgatan 22 (Kistagången 16)"
      , "City": "Stockholm" },
    "Number of lang instances": "2",
    "Abstract[eng ]": €€€€
    Today there are free online services that can be used to store files of arbitrary types and sizes, such as Google Drive. However, these services are often limited by a certain total storage size. The goal of this thesis is to create a filesystem that can store arbitrary amount and types of data, i.e. without any real limit to the total storage size. This is to be achieved by taking advantage of online webpages, such as Twitter, where text and files can be posted on free accounts with no apparent limit on storage size. The aim is to have a filesystem that behaves similar to any other filesystem but where the actual data is stored for free on various websites. €€€€,
    "Keywords[eng ]": €€€€
    Filesystem, Fejk FileSystem, Cloud-based filesystem, Steganograhpic filesystem €€€€,
    "Abstract[swe ]": €€€€
    Sammanfattning på svenska €€€€,
    "Keywords[swe ]": €€€€
    Filsystem, Fejk FileSystem, Molnbaserat filsystem, Steganografiskt filsystem €€€€,
  }
}
```